

Analysis of Yelp Business Intelligence Data

We will analyze a subset of Yelp's business, reviews and user data. This dataset comes to us from [Kaggle](#) although we have taken steps to pull this data into a public s3 bucket: `s3://sta9760-yelpdataset/yelp-light/*business.json`

Installation and Initial Setup

Begin by installing the necessary libraries that you may need to conduct your analysis. At the very least, you must install `pandas` and `matplotlib`

```
In [1]: sc.install_pypi_package("matplotlib==3.2.1")
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("seaborn==0.11.0")
sc.install_pypi_package("scipy==1.5.4")
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
5	application_1606229628628_0006	pyspark	idle	Link	Link	✓

SparkSession available as 'spark'.

Collecting matplotlib==3.2.1

Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl

Collecting python-dateutil>=2.1 (from matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/d4/70/d60450c3dd48ef87586924207ae8907090de0b306af2bce5d134d78615cb/python_dateutil-2.8.1-py2.py3-none-any.whl

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/8a/bb/488841f56197b13700afd5658fc279a2025a39e22449b7cf29864669b15d/pyparsing-2.4.7-py2.py3-none-any.whl>

Collecting cycler>=0.10 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/f7/d2/e07d3ebb2bd7af696440ce7e754c59dd546ffe1bbe732c8ab68b9c834e61/cyclor-0.10.0-py2.py3-none-any.whl>

Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)

Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)

Using cached https://files.pythonhosted.org/packages/d2/46/231de802ade4225b76b96cffe419cf3ce52bbe92e3b092cf12db7d11c207/kiwisolver-1.3.1-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)

Installing collected packages: python-dateutil, pyparsing, cycler, kiwisolver, matplotlib

Successfully installed cycler-0.10.0 kiwisolver-1.3.1 matplotlib-3.2.1 pyparsing-2.4.7 python-dateutil-2.8.1

Collecting pandas==1.0.3

Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1cdc0163743f55aeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from

```
python-dateutil>=2.6.1->pandas==1.0.3)
Installing collected packages: pandas
Successfully installed pandas-1.0.3
```

```
Collecting seaborn==0.11.0
```

```
Using cached https://files.pythonhosted.org/packages/bc/45/5118a05b0d61173e6eb12bc5804f0fbb6f196adb0a20e0b16efc2b8e98be/seaborn-0.11.0-py3-none-any.whl
```

```
Requirement already satisfied: numpy>=1.15 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.11.0)
```

```
Collecting scipy>=1.0 (from seaborn==0.11.0)
```

```
Using cached https://files.pythonhosted.org/packages/dc/7e/8f6a79b102ca1ea928bae8998b05bf5dc24a90571db13cd119f275ba6252/scipy-1.5.4-cp37-cp37m-manylinux1_x86_64.whl
```

```
Requirement already satisfied: matplotlib>=2.2 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from seaborn==0.11.0)
```

```
Requirement already satisfied: pandas>=0.23 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from seaborn==0.11.0)
```

```
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.0)
```

```
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.0)
```

```
Requirement already satisfied: cycler>=0.10 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.0)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages (from matplotlib>=2.2->seaborn==0.11.0)
```

```
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.23->seaborn==0.11.0)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib>=2.2->seaborn==0.11.0)
```

```
Installing collected packages: scipy, seaborn
```

```
Successfully installed scipy-1.5.4 seaborn-0.11.0
```

```
Requirement already satisfied: scipy==1.5.4 in /mnt/tmp/1606243180516-0/lib/python3.7/site-packages
```

```
Requirement already satisfied: numpy>=1.14.5 in /usr/local/lib64/python3.7/site-packages (from scipy==1.5.4)
```

Importing

Now, import the installed packages from the previous block below.

```
In [50]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
```

Loading Data

We are finally ready to load data. Using `spark` load the data from S3 into a `dataframe` object that we can manipulate further down in our analysis.

```
In [3]: business = spark.read.json('s3://sta9760-project2-yelpdataset/yelp_academic_dataset_bus
```

Overview of Data

Display the number of rows and columns in our dataset.

```
In [4]: num_rows = business.count()
num_cols = len(business.columns)
print('Columns:', num_cols, "| Rows:", num_rows)
```

Columns: 14 | Rows: 209393

Display the DataFrame schema below.

```
In [5]: business.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
|   |-- GoodForDancing: string (nullable = true)
|   |-- GoodForKids: string (nullable = true)
|   |-- GoodForMeal: string (nullable = true)
|   |-- HairSpecializesIn: string (nullable = true)
|   |-- HappyHour: string (nullable = true)
|   |-- HasTV: string (nullable = true)
|   |-- Music: string (nullable = true)
|   |-- NoiseLevel: string (nullable = true)
|   |-- Open24Hours: string (nullable = true)
|   |-- OutdoorSeating: string (nullable = true)
|   |-- RestaurantsAttire: string (nullable = true)
|   |-- RestaurantsCounterService: string (nullable = true)
|   |-- RestaurantsDelivery: string (nullable = true)
|   |-- RestaurantsGoodForGroups: string (nullable = true)
|   |-- RestaurantsPriceRange2: string (nullable = true)
|   |-- RestaurantsReservations: string (nullable = true)
|   |-- RestaurantsTableService: string (nullable = true)
|   |-- RestaurantsTakeOut: string (nullable = true)
|   |-- Smoking: string (nullable = true)
|   |-- WheelchairAccessible: string (nullable = true)
|   |-- WiFi: string (nullable = true)
|-- business_id: string (nullable = true)
|-- categories: string (nullable = true)
|-- city: string (nullable = true)
|-- hours: struct (nullable = true)
|   |-- Friday: string (nullable = true)
|   |-- Monday: string (nullable = true)
|   |-- Saturday: string (nullable = true)
|   |-- Sunday: string (nullable = true)
|   |-- Thursday: string (nullable = true)
```

```

|         |-- Tuesday: string (nullable = true)
|         |-- Wednesday: string (nullable = true)
|-- is_open: long (nullable = true)
|-- latitude: double (nullable = true)
|-- longitude: double (nullable = true)
|-- name: string (nullable = true)
|-- postal_code: string (nullable = true)
|-- review_count: long (nullable = true)
|-- stars: double (nullable = true)
|-- state: string (nullable = true)

```

Display the first 5 rows with the following columns:

- business_id
- name
- city
- state
- categories

```
In [6]: business.select("business_id", "name", "city", "state", "categories").show(5)
```

```

+-----+-----+-----+-----+-----+
| business_id | name | city | state | categories |
+-----+-----+-----+-----+-----+
| f9NumwFMBDn751xgF... | The Range At Lake... | Cornelius | NC | Active Life, Gun/... |
| YzvJg0SayhoZgCljU... | Carlos Santo, NMD | Scottsdale | AZ | Health & Medical,... |
| XNoUzKckATkOD1hP6... | Felinus | Montreal | QC | Pets, Pet Service... |
| 60AZjbxqM5o129BuH... | Nevada House of Hose | North Las Vegas | NV | Hardware Stores, ... |
| 51M2Kk903DFYI6gnB... | USE MY GUY SERVIC... | Mesa | AZ | Home Services, Pl... |
+-----+-----+-----+-----+-----+

```

only showing top 5 rows

Analyzing Categories

Let's now answer this question: **how many unique categories are represented in this dataset?**

Essentially, we have the categories per business as a list - this is useful to quickly see what each business might be represented as but it is difficult to easily answer questions such as:

- How many businesses are categorized as Active Life , for instance
- What are the top 20 most popular categories available?

Association Table

We need to "break out" these categories from the business ids? One common approach to take is to build an association table mapping a single business id multiple times to each distinct category.

For instance, given the following:

business_id	categories
abcd123	a,b,c

We would like to derive something like:

business_id	category
abcd123	a
abcd123	b
abcd123	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Implement the code necessary to derive the table described from your original yelp dataframe.

```
In [7]: from pyspark.sql.functions import explode, split
```

```
In [8]: business_categories = business.select("business_id", "categories")
        business_categories_exploded = business_categories.withColumn('categories', explode(split(business_categories.categories, ',')))
```

Display the first 5 rows of your association table below.

```
In [9]: business_categories_exploded.show(5)
```

```
+-----+-----+
| business_id | categories |
+-----+-----+
| f9NumwFMBDn751xgF... | Active Life |
| f9NumwFMBDn751xgF... | Gun/Rifle Ranges |
| f9NumwFMBDn751xgF... | Guns & Ammo |
| f9NumwFMBDn751xgF... | Shopping |
| YzvJg0SayhoZgCljU... | Health & Medical |
+-----+-----+
only showing top 5 rows
```

Total Unique Categories

Finally, we are ready to answer the question: **what is the total number of unique categories available?**

Below, implement the code necessary to calculate this figure.

```
In [10]: business_categories_exploded.select('categories').distinct().count()
```

1336

Top Categories By Business

Now let's find the top categories in this dataset by rolling up categories.

Counts of Businesses / Category

So now, let's unroll our distinct count a bit and display the per count value of businesses per category.

The expected output should be:

category	count
a	15
b	2
c	45

Or something to that effect.

```
In [11]: business_categories_exploded.groupby('categories').count().show()
```

```
+-----+-----+
|      categories|count|
+-----+-----+
| Paddleboarding|  36|
| Dermatologists| 341|
| Aerial Tours  |  28|
| Hobby Shops   | 828|
| Bubble Tea    | 720|
| Embassy       |  13|
| Tanning       | 938|
| Handyman      | 682|
| Aerial Fitness|  29|
| Falafel       | 159|
| Outlet Stores | 399|
| Summer Camps  | 318|
| Clothing Rental|  55|
| Sporting Goods|2311|
| Cooking Schools| 118|
| College Counseling| 15|
| Lactation Services| 50|
| Ski & Snowboard S...| 50|
| Museums       | 359|
| Doulas        |  45|
+-----+-----+
only showing top 20 rows
```

Bar Chart of Top Categories

With this data available, let us now build a barchart of the top 20 categories.

HINT: don't forget about the matplotlib magic!

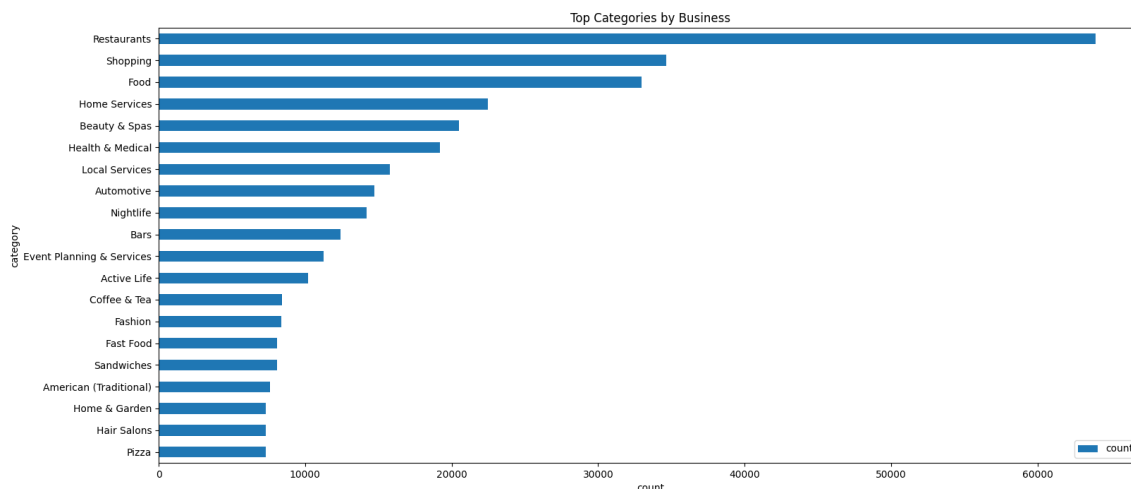
```
%matplotlib plt
```

```
In [12]: barchart_business = business_categories_exploded.groupby('categories').count().orderBy(
pdf = barchart_business.limit(20).toPandas())
```

```
In [13]: pdf = pdf.sort_values('count', ascending=True)
pdf.plot(kind='barh', x='categories', y='count', figsize=(18, 8))
```

```
plt.title('Top Categories by Business')
plt.xlabel('count')
plt.ylabel('category')
```

```
%matplotlib plt
```



Do Yelp Reviews Skew Negative?

Oftentimes, it is said that the only people who write a written review are those who are extremely *dissatisfied* or extremely *satisfied* with the service received.

How true is this really? Let's try and answer this question.

Loading User Data

Begin by loading the user data set from S3 and printing schema to determine what data is available.

```
In [14]: review = spark.read.json('s3://sta9760-project2-yelpdataset/yelp_academic_dataset_revie
review.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

Let's begin by listing the `business_id` and `stars` columns together for the user reviews data.

```
In [15]: review.select('business_id', 'stars').show(5)
```

```

+-----+-----+
|      business_id|stars|
+-----+-----+
|-MhfebM0QIsKt87iD...| 2.0|
|lbrU8StCq3yDfr-QM...| 1.0|
|HQ128KMwrEKHqhFrr...| 5.0|
|5Jx1ZaqCnk1MnbgRi...| 1.0|
|IS4cv902ykd8wj1TR...| 4.0|
+-----+-----+
only showing top 5 rows

```

Now, let's aggregate along the `stars` column to get a resultant dataframe that displays *average stars* per business as accumulated by users who **took the time to submit a written review**.

```

In [16]: from pyspark.sql.functions import length
          from pyspark.sql.functions import col

          written_reviews = review.select('business_id', 'user_id', 'stars', 'text').where(length
          avg_stars = written_reviews.groupby("business_id").avg("stars")
          avg_stars.show(5)

```

```

+-----+-----+
|      business_id|      avg(stars)|
+-----+-----+
|VHsNB3pdGVcRgs6C3...| 3.411764705882353|
|RMjCnixEY5i12Ciqn...| 3.5316455696202533|
|ipFreSFhjClfNETuM...| 2.6|
|dLDMU8bOLnkDTmPUr...| 4.942857142857143|
|Qm2datcYBPXrPATVG...| 4.352941176470588|
+-----+-----+
only showing top 5 rows

```

Now the fun part - let's join our two dataframes (reviews and business data) by `business_id`.

```

In [17]: business = business.join(avg_stars, 'business_id', 'inner')

```

Let's see a few of these:

```

In [18]: business.select('avg(stars)', 'stars', 'name', 'city', 'state').show(5)

```

```

+-----+-----+-----+-----+-----+
|      avg(stars)|stars|      name|      city|state|
+-----+-----+-----+-----+-----+
|4.11784140969163| 4.0|Delmonico Steakhouse|Las Vegas|NV|
|                | 4.5|Mr. Pancho Mexica...|Mesa|AZ|
|                | 3.75|Maricopa County D...|Phoenix|AZ|
|                | 4.0|Double Play Sport...|Las Vegas|NV|
|2.6875| 2.5|Impressions Dental|Chandler|AZ|
+-----+-----+-----+-----+-----+
only showing top 5 rows

```

Compute a new dataframe that calculates what we will call the *skew* (for lack of a better word) between the avg stars accumulated from written reviews and the *actual* star rating of a business (ie: the average of stars given by reviewers who wrote an actual review **and** reviewers who just provided a star rating).

The formula you can use is something like:

$$(\text{row}['\text{avg}(\text{stars})'] - \text{row}['\text{stars}']) / \text{row}['\text{stars}']$$

If the **skew** is negative, we can interpret that to be: reviewers who left a written response were more dissatisfied than normal. If **skew** is positive, we can interpret that to be: reviewers who left a written response were more satisfied than normal.

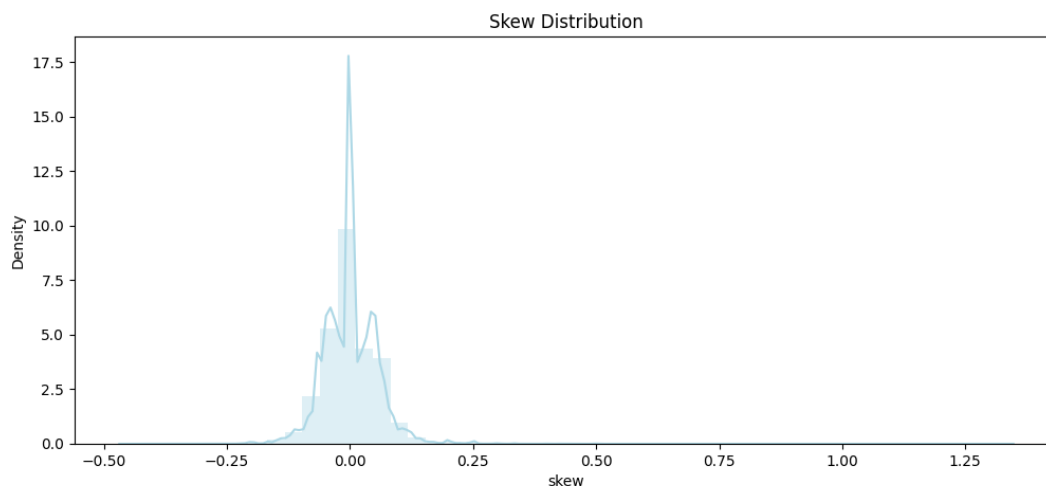
```
In [19]: skew_df = business.select('avg(stars)', 'stars', 'name', 'city', 'state').withColumn('skew', (row['avg(stars)'] - row['stars']) / row['stars']).show(10)
```

avg(stars)	stars	name	city	state	skew
4.11784140969163	4.0	Delmonico Steakhouse	Las Vegas	NV	0.029460352422907565
4.5	4.5	Mr. Pancho Mexica...	Mesa	AZ	0.0
3.75	4.0	Maricopa County D...	Phoenix	AZ	-0.0625
4.0	4.0	Double Play Sport...	Las Vegas	NV	0.0
2.6875	2.5	Impressions Dental	Chandler	AZ	0.075
4.976744186046512	5.0	Kidz Cuts By Lori	Henderson	NV	-0.00465116279069...
3.8107142857142855	4.0	Río Mirage Café y...	El Mirage	AZ	-0.04732142857142...
3.7941176470588234	4.0	Steep & Brew West	Madison	WI	-0.05147058823529416
1.4762931034482758	1.5	Showtime Tours	Las Vegas	NV	-0.01580459770114...
2.0	2.0	August Moon Chine...	Woodbridge	ON	0.0

only showing top 10 rows

And finally, graph it!

```
In [20]: pd_skew = skew_df.toPandas()
fig = plt.figure(figsize=(12,5))
ax = sns.distplot(pd_skew['skew'], hist=True, kde=True, color = 'lightblue').set_title('Skew Distribution')
plt.show()
%matplotlib plt
```



So, do Yelp (written) Reviews skew negative? Does this analysis actually prove anything? Expound on implications / interpretations of this graph.

From this distribution, we can see that the Yelp reviews are skewed to the left. If the skew number is 0, it means that the reviews are neutral and not skewed. However, there are many more reviews with

negative skew numbers. Thus, reviewers who left a written response were more dissatisfied than normal

Should the Elite be Trusted? (Or, some other analysis of your choice)

For the final portion - you have a choice:

- Try and analyze some interesting dimension to this data. The **ONLY** requirement is that you must use the **Users** dataset and join on either the **business*** or **reviews**** dataset
- Or, you may try and answer the question posed: how accurate or close are the ratings of an "elite" user (check Users table schema) vs the actual business rating.

Feel free to use any and all methodologies at your disposal - only requirement is you must render one visualization in your analysis

```
In [21]: user = spark.read.json('s3://sta9760-project2-yelpdataset/yelp_academic_dataset_user.js')
user.printSchema()
```

```
root
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)
```

```
In [22]: user.select('user_id', 'average_stars', 'elite').show(5, truncate=False)
```

user_id	average_stars	elite
ntlrvfPzc8eglqvK92iDIAw	3.57	
FOBRP1BHa3WPHFB5qYD1Vg	3.84	2008, 2009, 2010, 2011, 2012, 2013
zZUnPeh2hEp0WydbAZE0Og	3.44	2010
QaELAmRcDc5TfJEy1aaP8g	3.08	2009
xvu8G900tezTzbbfqmTKvA	4.37	2009, 2010, 2011, 2012, 2014, 2015, 2016, 2017, 2018

only showing top 5 rows

Create the elite user subset from the total user dataset

```
In [23]: import pyspark.sql.functions as F
df = user.select('user_id', 'average_stars', 'elite').withColumn("length_of_elite", F.length_of_elite)
df.show(truncate=False)
df.count()
```

```
+-----+-----+-----+-----+
|user_id|average_stars|elite|length_of_elite|
+-----+-----+-----+-----+
|ntlvfPzc8eglqvK92iDIAw|3.57| |0|
|FOBRP1BHa3WPHFB5qYD1Vg|3.84|2008,2009,2010,2011,2012,2013|29|
|zZUnPeh2hEp0WydbAZE00g|3.44|2010|4|
|QaELAmRcDc5TfJEy1aaP8g|3.08|2009|4|
|xvu8G900tezTzbbfqmTKvA|4.37|2009,2010,2011,2012,2014,2015,2016,2017,2018|44|
|z5_82komKV3mI4ASGe2-FQ|2.88|2007|4|
|ttumcu6hWshk_EJvWrduDg|4.0| |0|
|f4_MRNHvN-yRn7EA8YWRxg|3.63|2011,2012,2013,2014,2015,2016,2017,2018|39|
|UYACF30806j2mfbB5vdmJA|3.75| |0|
|QG13XBbgHwydzThRBGJtyw|4.1|2008,2009|9|
|f6YuZP6iennHFVlnFJOXLQ|3.8| |0|
|I_6wY8_RsewziNnKhGZg4g|3.63|2010,2011,2012,2013,2014,2015,2016,2017|39|
|q-v8e1VPvKz0KvK69QSj1Q|3.37|2011,2012,2013,2014,2015,2016,2017,2018|39|
|HwPGLzF_uXB3MF8bc5u5dg|4.5| |0|
|y4UuVowA9i3zj2hHyRMfHw|4.17| |0|
|1WBxJ2r3A2QYfRSEzgcmkQ|3.82|2010,2011,2012,2013,2014,2015,2016|34|
|-TT5e-YQU9xLb1JAGCGkQw|3.91|2010,2011,2012,2013|19|
|6bbHSJ0PrgSxh7e5nigKMw|2.21| |0|
|4VmuXuSRhv5UxYUy3tMpiQ|3.88|2012,2013|9|
|pVU2DdtBFppBAX5G5t3rcw|3.79| |0|
+-----+-----+-----+-----+
only showing top 20 rows
```

1968703

In total, there are 1968703 users in this dataset.

```
In [24]: elite = user.select('user_id', 'elite').where(length(col("elite")) > 0)
```

```
elite.count()
```

75961

There is a total of 75961 elite users throughout the years in this dataset.

Join elite user dataset with review dataset on user_id to find out the reviews they have submitted:

```
In [25]: elite_reviews = elite.join(review, 'user_id', 'inner')
         elite_reviews.select('business_id', 'user_id', 'review_id', 'stars').show(10)
```

business_id	user_id	review_id	stars
-8F04F54iDT6VgWPC...	1Du159QEe-Q-70QHT...	Kg5ncegiJ3utWRxDt...	4.0
p200k46G_A000nCw1...	3pMczoCBOSKBcqMhV...	nBNDv9j_tiPPo5MMe...	5.0
jyFoxS8MofdpkAAK6...	j044Apni7iJZVVK4H...	EI6L-L0Dcj6HAUaB0...	1.0
ewty6EB70nwPJSUkA...	R078oDy7vbEc0JU8a...	Ryohf9HJcpk2C49vf...	4.0
0M3KCmdY-_x1Iu5vE...	TFxeEvpjMNQ3AWL49...	Lc0Tj-Me2Jwu_V9au...	5.0
-h0o-BilkKaCa7HX9...	F11oTs6usaCfyjLnY...	GORTMUfkTtGViv4ap...	5.0
DEtOIjhV0MMWZ8fD8...	R078oDy7vbEc0JU8a...	aULkXMFrsMvctmJ5Q...	5.0
Jt28TYWanzKrJYYr0...	LEr8vS6PRymCg-SJH...	JfNrW6b2mgcynJ3w2...	2.0
NFm869_w6cvVaWaNP...	M7vDDzoPNQDN2FdTc...	OcaQAZf1KbxKLS2rT...	5.0
Da6eZFThE9xanUAGN...	Ania9MCwET-TBzVjV...	wuNIHeqK_pjpE4Hrp...	4.0

only showing top 10 rows

Calculate sum of elite user ratings each business received:

```
In [26]: ereview_count_by_business = elite_reviews.select('business_id', 'user_id', 'review_id',
         ereview_count_by_business = ereview_count_by_business.withColumnRenamed("count", "erevie
         ereview_count_by_business.show(5)
```

business_id	ereview_count
VHsNB3pdGVcRgs6C3...	25
-I06hkMFrX0KBqu61...	1
RMjCnixEY5i12Ciqn...	26
ipFreSFhjClfNETuM...	17
Qm2datcYBPXrPATVG...	3

only showing top 5 rows

Calculate the average elite rating each business received:

```
In [27]: avg_reviews_by_elite = elite_reviews.select('business_id', 'user_id', 'review_id', 'sta
         avg_reviews_by_elite = avg_reviews_by_elite.withColumnRenamed("avg(stars)", "avg_e_stars
         avg_reviews_by_elite = avg_reviews_by_elite.join(ereview_count_by_business, 'business_i
         avg_reviews_by_elite.show(5)
```

business_id	avg_e_stars	ereview_count
--9e10NYQuAa-CB_R...	4.1916058394160585	548

```
--phjqoPSPa8sLmUV...| 4.0| 4|
--q7kSBRb0vWC8lSk...| 4.0| 1|
-0Z000Vm2ADchyt1E...| 5.0| 8|
-1VaIJza42Hjev6uk...| 3.793103448275862| 29|
+-----+-----+
only showing top 5 rows
```

Join the average elite stars dataset with the business dataset:

```
In [28]: avg_reviews_by_elite = avg_reviews_by_elite.join(business, 'business_id', 'inner')
avg_reviews = avg_reviews_by_elite.select('business_id', 'stars', 'avg_e_stars', 'erevi
avg_reviews.show(5)
```

```
+-----+-----+-----+-----+
| business_id|stars| avg_e_stars|ereview_count|
+-----+-----+-----+-----+
|--9e10NYQuAa-CB_R...| 4.0|4.1916058394160585| 548|
|--phjqoPSPa8sLmUV...| 4.0| 4.0| 4|
--q7kSBRb0vWC8lSk...| 4.0| 4.0| 1|
-0Z000Vm2ADchyt1E...| 5.0| 5.0| 8|
-1VaIJza42Hjev6uk...| 4.0| 3.793103448275862| 29|
+-----+-----+-----+-----+
only showing top 5 rows
```

Calculate sum of all user ratings each business received:

```
In [29]: allreview_count_by_business = review.groupby('business_id').count()
allreview_count_by_business = allreview_count_by_business.withColumnRenamed("count", "al
allreview_count_by_business.show(5)
```

```
+-----+-----+
| business_id|allreview_count|
+-----+-----+
|VHsNB3pdGVcRgs6C3...| 136|
|RMjCnixEY5i12Ciqn...| 79|
|ipFreSFhjClfNETuM...| 80|
|dLDMU8b0LnkDTmPUr...| 35|
|Qm2datcYBPXrPATVG...| 17|
+-----+-----+
only showing top 5 rows
```

Joining all rating data together:

```
In [30]: avg_reviews = avg_reviews.join(allreview_count_by_business, 'business_id', 'inner')
avg_reviews.show(5)
```

```
+-----+-----+-----+-----+-----+
| business_id|stars| avg_e_stars|ereview_count|allreview_count|
+-----+-----+-----+-----+-----+
|--9e10NYQuAa-CB_R...| 4.0|4.1916058394160585| 548| 1816|
|--phjqoPSPa8sLmUV...| 4.0| 4.0| 4| 12|
--q7kSBRb0vWC8lSk...| 4.0| 4.0| 1| 7|
-0Z000Vm2ADchyt1E...| 5.0| 5.0| 8| 86|
-1VaIJza42Hjev6uk...| 4.0| 3.793103448275862| 29| 280|
+-----+-----+-----+-----+-----+
only showing top 5 rows
```

Calculate the percentage of elite ratings each business received:

```
In [31]: avg_reviews = avg_reviews.withColumn('elite_percentage', (col('ereview_count') / col('allreview_count')) * 100)
avg_reviews = avg_reviews.withColumnRenamed("stars","avg_stars")
avg_reviews.show(5)
```

```
+-----+-----+-----+-----+-----+-----+
| business_id | avg_stars | avg_e_stars | ereview_count | allreview_count | elite_percentage |
+-----+-----+-----+-----+-----+-----+
|--9e10NYQuAa-CB_R... | 4.0 | 4.1916058394160585 | 548 | 1816 | 30.1762
11453744493 |
|--phjqoPSPa8sLmUV... | 4.0 | 4.0 | 4 | 12 | 33.333
33333333333 |
|--q7kSBRb0vWC8lSk... | 4.0 | 4.0 | 1 | 7 | 14.2857
14285714285 |
|-0Z000Vm2ADchyt1E... | 5.0 | 5.0 | 8 | 86 | 9.302
32558139535 |
|-1VaIJza42Hjev6uk... | 4.0 | 3.793103448275862 | 29 | 280 | 10.3571
42857142858 |
+-----+-----+-----+-----+-----+-----+
```

only showing top 5 rows

Let's take a look at the top 25 businesses with the highest number of reviews

```
In [32]: avg_reviews.orderBy('allreview_count', ascending=False).show(25)
```

```
+-----+-----+-----+-----+-----+-----+
| business_id | avg_stars | avg_e_stars | ereview_count | allreview_count | elite_percentage |
+-----+-----+-----+-----+-----+-----+
| RESDUcs7fIiihp38-... | 4.0 | 4.060636515912898 | 2985 | 10417 | 28.6550
83037342806 |
| 4JNXUYy8wbbaaDmk3B... | 4.0 | 4.13179992698065 | 2739 | 9536 | 28.7227
34899328863 |
| K7lWdNUhCbcnEvI0N... | 3.5 | 3.8233124308373294 | 2711 | 7594 | 35.699
23623913616 |
| f4x1YBxkLrZg652xt... | 4.0 | 3.901702361339923 | 1821 | 6859 | 26.549
05962968363 |
| cYwJA2A6I12KNkm2r... | 4.0 | 3.938221317040054 | 1473 | 5586 | 26.369
49516648765 |
| DkYS3arLOhA8si5uU... | 4.5 | 4.27808988764045 | 2492 | 5370 | 46.405
95903165735 |
| faPVqws-x-5k2CQKD... | 4.5 | 4.403796376186367 | 1159 | 4979 | 23.2777
66619803174 |
| 5LNZ67Yw9RD6nf4_U... | 4.0 | 4.263126131563066 | 1657 | 4973 | 33.319
92760908908 |
| 2weQS-Rno0Bhb1KsH... | 3.5 | 3.819277108433735 | 1826 | 4953 | 36.866
54552796285 |
| iCQpiavjjPzJ5_3gP... | 4.0 | 4.105722599418041 | 2062 | 4882 | 42.236
78820155674 |
| AV6weBrZFbFRGcbc... | 2.5 | 2.957685664939551 | 1158 | 4819 | 24.0298
81718198798 |
| vHz2RLtFUMVRPFmd7... | 4.5 | 4.394230769230769 | 728 | 4801 | 15.1635
07602582795 |
+-----+-----+-----+-----+-----+-----+
```

SMPbvZLSMMb7KU76Y...	3.5	3.845360824742268	1455	4749	30.638
02905874921					
ujHiaprwCQ5ewziu0...	3.5	3.616777308388654	1657	4731	35.024
30775734517					
El4FC8jcawUVgw_0E...	3.0	3.4377713458755426	1382	4589	30.1154
93571584224					
QXV3L_QFGj8r6nWX2...	4.5	4.193423597678917	517	4357	11.8659
62818453065					
rcaPajgK0JC2vo_13...	4.0	4.120710059171597	1690	4305	39.2566
78281068524					
JDZ6_yycNQFTpUZzL...	4.5	4.154639175257732	388	4225	9.1834
31952662723					
OETH78qcgDltvHULo...	4.0	4.213844252163164	809	4217	19.1842
54209153426					
3kdSl5mo9dWC4clrQ...	4.5	4.297560975609756	615	4125	14.9090
90909090908					
YJ81jUhLsz6CtT_20...	3.5	3.62015503875969	774	4119	18.7909
68681718866					
KskYqH1Bi7Z_61pH6...	4.0	4.215588723051409	1206	4119	29.2
78951201748					
RwMLuOkImBIqqYj4S...	4.0	4.294117647058823	1020	4088	24.9510
76320939332					
FaHADZARwnY4yvlvp...	3.5	3.411930673115679	2481	4064	61.048
22834645669					
u_vPjx925UPEG9DFO...	2.5	3.091928251121076	892	4024	22.1669
98011928428					

```
+-----+-----+-----+-----+-----+-----+
+-----+
```

only showing top 25 rows

```
In [33]: avg_reviews.summary().show()
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|summary|          business_id|          avg_stars|          avg_e_stars|          ereview_count|
|allreview_count| elite_percentage|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|  count|          148225|          148225|          148225|          148225|
148225|          148225|
|  mean|          null|3.571276775172879| 3.781229949617148|11.84906054983977|51.
138910440209145|29.071736936940418|
| stddev|          null|0.904201241413331|0.9734324285178444|41.04359275261798| 14
9.5476463476886| 19.78786142447515|
|  min|--1UhmGODdWsrMast...|          1.0|          1.0|          1|
3|0.5208333333333333|
|  25%|          null|          3.0| 3.232142857142857|          1|
6|14.285714285714285|
|  50%|          null|          3.5|          4.0|          3|
15|          25.0|
|  75%|          null|          4.0|          4.5|          9|
43| 38.83495145631068|
|  max|zzaIBwimxVej4tY6...|          5.0|          5.0|          2985|
10417|          100.0|
+-----+-----+-----+-----+-----+-----+
+-----+
```

- There are many businesses with very few reviews. We are going to filter them out and keep only businesses that have at least 25 total reviews.
- We also want to filter out businesses that have less than 5 elite reviews and also businesses where elite reviews account for less than 5% of total reviews.

- This way, we have quality data points to calculate the businesses' average reviews and average elite reviews.

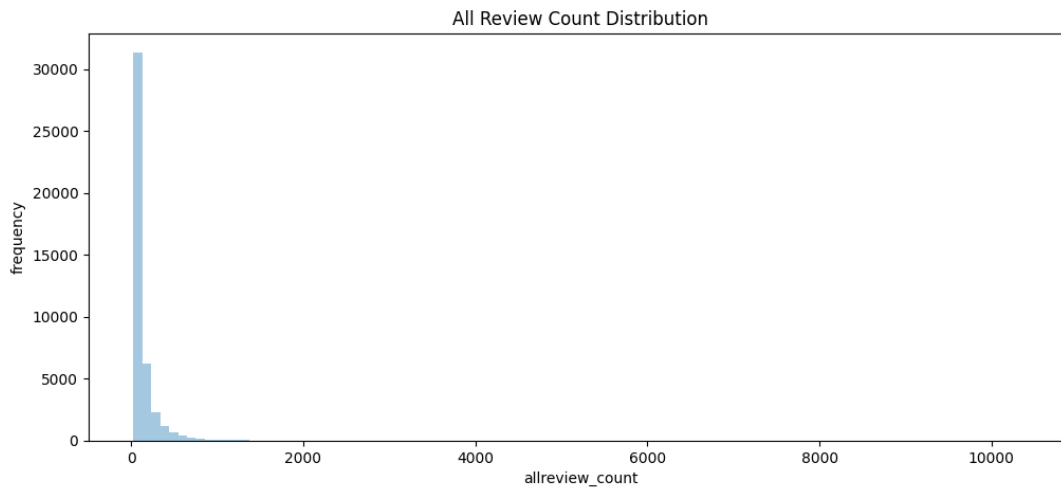
```
In [34]: avg_reviews = avg_reviews.where(col("allreview_count") >= 25)
avg_reviews = avg_reviews.where(col("ereview_count") >= 5)
avg_reviews = avg_reviews.where(col("elite_percentage") >= 5)
```

```
In [35]: avg_reviews.summary().show()
```

summary	business_id	avg_stars	avg_e_stars	ereview_count
allreview_count	elite_percentage			
count	43291	43291	43291	43291
mean	3.6322561271395903	3.781319959083377	33.84530271880991	139.58448638285094
stddev	0.7020745673904851	0.5875925928795263	71.2158803696474	253.57361456349582
min	1.0	1.0	5	25
25%	3.0	3.4166666666666665	9	42
50%	3.5	3.8260869565217392	17	72
75%	4.0	4.186440677966102	33	140
max	5.0	5.0	2985	10417

Our dataset looks a lot better after some filtering and cleaning.

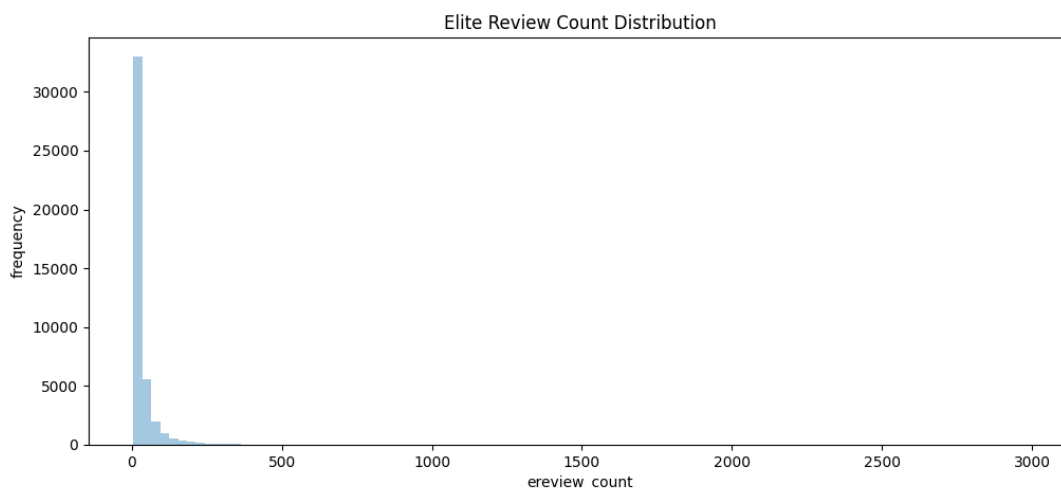
```
In [36]: pd_avg_reviews = avg_reviews.toPandas()
fig = plt.figure(figsize=(12,5))
ax = sns.distplot(pd_avg_reviews['allreview_count'], hist=True, bins=100, kde=False).se
plt.xlabel("allreview_count")
plt.ylabel("frequency")
plt.show()
%matplotlib plt
```

```
In [37]: pd_avg_reviews['allreview_count'].mode()
```

```
0    25
dtype: int64
```

```
In [38]: fig = plt.figure(figsize=(12,5))
ax = sns.distplot(pd_avg_reviews['ereview_count'], hist=True, bins=100, kde=False).set_
plt.xlabel("ereview_count")
plt.ylabel("frequency")
plt.show()
%matplotlib plt
```



Let's take a look at the difference between average elite rating and average rating:

```
In [39]: avg_reviews = avg_reviews.withColumn('delta', (col('avg_e_stars') - col('avg_stars')))
avg_reviews.show(10)
```

```
+-----+-----+-----+-----+-----+
| business_id | avg_stars | avg_e_stars | ereview_count | allreview_count | elite |
|_percentage|_delta|
+-----+-----+-----+-----+-----+-----+
```

```

-----+-----+
|--9e10NYQuAa-CB_R...|      4.0|4.1916058394160585|      548|      1816|30.1762
11453744493|0.19160583941605847|
|-0Z000Vm2ADchyt1E...|      5.0|      5.0|      8|      86|  9.302
32558139535|      0.0|
|-1VaIJza42Hjev6uk...|      4.0| 3.793103448275862|      29|      280|10.3571
42857142858|-0.2068965517241379|
|-2Arz8twKJmxHMS3S...|      4.0|      3.7|      10|      34|29.4117
64705882355|-0.29999999999999998|
|-2ToCaDFpTNmmg3QF...|      1.5|      2.109375|      64|      464|13.7931
03448275861|      0.609375|
|-2wh_ZsD2n5xFYgzp...|      4.0|3.7058823529411766|      17|      34|
50.0|-0.2941176470588234|
|-5NXoZeGBdx3Bdk70...|      4.0|      3.4|      10|      76|13.1578
94736842104|-0.6000000000000001|
|-5__awbuGMHAK6cZg...|      3.0|2.4285714285714284|      7|      78|  8.9743
58974358974|-0.5714285714285716|
|-ArzV0ksIBWmtM1ey...|      4.5| 4.666666666666667|      12|      57|21.0526
31578947366|0.16666666666666696|
|-BbnAc9YE06pJvJGE...|      4.0|4.2342342342342345|      111|      270| 41.111
11111111111| 0.2342342342342345|
-----+-----+

```

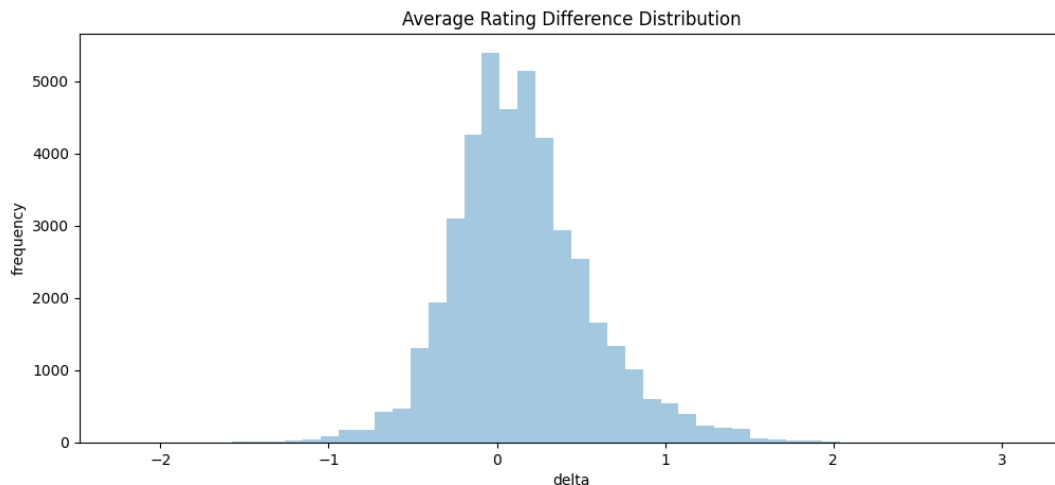
only showing top 10 rows

Now, let's look at the distribution of this delta:

```

In [40]: pd_avg_reviews = avg_reviews.toPandas()
fig = plt.figure(figsize=(12,5))
ax = sns.distplot(pd_avg_reviews['delta'], hist=True, kde=False).set_title('Average Rat
plt.xlabel("delta")
plt.ylabel("frequency")
plt.show()
%matplotlib plt

```



We see that the difference between average elite reviews and average total reviews is normally distributed. There might be some skewness to the right. We will first look at the summary of the distribution:

```

In [41]: avg_reviews.select('delta').summary().show()

```

```

-----+-----+

```

summary	delta
count	43291
mean	0.149063831943786
stddev	0.42068188793625066
min	-2.2142857142857144
25%	-0.11538461538461542
50%	0.11224489795918347
75%	0.375
max	3.0999999999999996

From an initial glance, we see that the mean is around 0.15 which suggests that elite users give a rating of 0.15 stars higher than the average of total users' ratings. The standard deviation is 0.42 which suggests that we should perform a hypothesis test in order to statistically reject the null hypothesis (average rating of elite users equals the average rating of all users).

Hypothesis Testing

1. Determine a null and alternate hypothesis:

- Null Hypothesis: Average rating of elite users and all users are the same
- Alternate Hypothesis: Average rating of elite users and all users are different

2. Sample size

From the summary above of the dataset, the sample size for our test is 43291. (N = 43291)

```
In [42]: N = 43291
```

3. Determine a confidence interval and degrees of freedom

We select $\alpha = 0.05$ - there is 95% confidence that the conclusion of this test will be valid.

Degree of freedom:

df = sample size (elite users) + sample size (all users) - 2 = 43291 + 43291 - 2

```
In [43]: degree_freedom = 2*avg_reviews.count() - 2
```

```
In [44]: print(degree_freedom)
```

86580

4. Calculate standard deviation

Calculate the variance

```
In [45]: from pyspark.sql.functions import var_samp
var_elite = avg_reviews.select(var_samp("avg_e_stars")).head()[0]
var_all = avg_reviews.select(var_samp("avg_stars")).head()[0]
```

```
In [46]: print(var_elite)
         print(var_all)
```

```
0.3452650552068847
0.4929086981765366
```

Calculate the standard deviation

```
In [47]: s = np.sqrt((var_elite + var_all)/2)
         print(s)
```

```
0.6473691965885546
```

5. Calculate the t-statistics

```
In [48]: # From our statistic summary of dataset
         mean_elite = 3.781319959083376
         mean_all = 3.6322561271395903
         t = (mean_elite - mean_all)/(s*np.sqrt(2/N))
         print(t)
```

```
33.876931165656266
```

6. Compare with critical value

p-value after comparison with t

```
In [49]: p = 1 - stats.t.cdf(t,df=degree_freedom)
         print("t = " + str(t))
         print("p = " + str(2*p))
```

```
t = 33.876931165656266
p = 0.0
```

After comparing the t statistic with the critical t value (using the stats module), we obtained a p value of 0.0. A p-value less than 0.05 (typically ≤ 0.05) is statistically significant. It indicates strong evidence against the null hypothesis, as there is less than a 5% probability the null is correct (and the results are random).

Therefore, we reject the null hypothesis, and accept the alternative hypothesis. Thus, it proves that the average rating of elite users and all users are different and that the elite reviewers should not be trusted.